

# Assignment 2

February 9, 2020

---

You are currently looking at **version 1.2** of this notebook. To download notebooks and datafiles, as well as get help on Jupyter notebooks in the Coursera platform, visit the [Jupyter Notebook FAQ](#) course resource.

---

## 1 Assignment 2 - Pandas Introduction

All questions are weighted the same in this assignment. ## Part 1 The following code loads the olympics dataset (olympics.csv), which was derived from the Wikipedia entry on [All Time Olympic Games Medals](#), and does some basic data cleaning.

The columns are organized as # of Summer games, Summer medals, # of Winter games, Winter medals, total # number of games, total # of medals. Use this dataset to answer the questions below.

```
In [1]: import pandas as pd
```

```
df = pd.read_csv('olympics.csv', index_col=0, skiprows=1)
```

```
for col in df.columns:
    if col[:2]=='01':
        df.rename(columns={col:'Gold'+col[4:]}, inplace=True)
    if col[:2]=='02':
        df.rename(columns={col:'Silver'+col[4:]}, inplace=True)
    if col[:2]=='03':
        df.rename(columns={col:'Bronze'+col[4:]}, inplace=True)
    if col[:1]==' ':
        df.rename(columns={col:'#'+col[1:]}, inplace=True)
```

```
names_ids = df.index.str.split('\s\(') # split the index by '('
```

```
df.index = names_ids.str[0] # the [0] element is the country name (new index)
```

```
df['ID'] = names_ids.str[1].str[:3] # the [1] element is the abbreviation or ID (take first 3 characters)
```

```
df = df.drop('Totals')
```

```
df.head()
```

```

Out[1]:
# Summer  Gold  Silver  Bronze  Total  # Winter  Gold.1  \
Afghanistan      13     0     0       2       2         0     0
Algeria           12     5     2       8      15         3     0
Argentina         23    18    24      28      70        18     0
Armenia            5     1     2       9      12         6     0
Australasia       2     3     4       5      12         0     0

Silver.1  Bronze.1  Total.1  # Games  Gold.2  Silver.2  Bronze.2  \
Afghanistan      0         0         0       13         0         0         2
Algeria           0         0         0       15         5         2         8
Argentina         0         0         0       41        18        24        28
Armenia            0         0         0       11         1         2         9
Australasia       0         0         0        2         3         4         5

Combined total  ID
Afghanistan         2  AFG
Algeria             15  ALG
Argentina           70  ARG
Armenia             12  ARM
Australasia         12  ANZ

```

### 1.0.1 Question 0 (Example)

What is the first country in df?

*This function should return a Series.*

```

In [2]: # You should write your whole answer within the function provided. The autograder will call
# this function and compare the return value against the correct solution value
def answer_zero():
    # This function returns the row for Afghanistan, which is a Series object. The assignment
    # question description will tell you the general format the autograder is expecting
    return df.iloc[0]

# You can examine what your function returns by calling it in the cell. If you have questions
# about the assignment formats, check out the discussion forums for any FAQs
answer_zero()

```

```

Out[2]: # Summer      13
Gold           0
Silver         0
Bronze         2
Total          2
# Winter       0
Gold.1         0
Silver.1       0
Bronze.1       0
Total.1        0
# Games        13

```

```

Gold.2          0
Silver.2        0
Bronze.2        2
Combined total   2
ID              AFG
Name: Afghanistan, dtype: object

```

### 1.0.2 Question 1

Which country has won the most gold medals in summer games?

*This function should return a single string value.*

```

In [4]: def answer_one():
        return df['Gold'].idxmax()

        answer_one()

```

```
Out[4]: 'United States'
```

### 1.0.3 Question 2

Which country had the biggest difference between their summer and winter gold medal counts?

*This function should return a single string value.*

```

In [6]: def answer_two():
        df['diff'] = df['Gold']-df['Gold.1']
        return df['diff'].idxmax()

        answer_two()

```

```
Out[6]: 'United States'
```

### 1.0.4 Question 3

Which country has the biggest difference between their summer gold medal counts and winter gold medal counts relative to their total gold medal count?

$$\frac{\text{Summer Gold} - \text{Winter Gold}}{\text{Total Gold}}$$

Only include countries that have won at least 1 gold in both summer and winter.

*This function should return a single string value.*

```

In [9]: def answer_three():
        df_nonzero = df[(df['Gold']>0) & (df['Gold.1']>0)]
        df_nonzero['diff_tot'] = (df_nonzero['Gold']-df_nonzero['Gold.1'])/df['Gold.2']
        return df_nonzero['diff_tot'].idxmax()

        answer_three()

```

```
Out[9]: 'Bulgaria'
```

### 1.0.5 Question 4

Write a function that creates a Series called "Points" which is a weighted value where each gold medal (Gold.2) counts for 3 points, silver medals (Silver.2) for 2 points, and bronze medals (Bronze.2) for 1 point. The function should return only the column (a Series object) which you created, with the country names as indices.

*This function should return a Series named Points of length 146*

```
In [12]: def answer_four():  
         df['Total Points'] = (df['Gold.2']*3)+(df['Silver.2']*2)+(df['Bronze.2'])  
         return df['Total Points']
```

```
answer_four()
```

```
Out[12]: Afghanistan      2  
         Algeria          27  
         Argentina       130  
         Armenia         16  
         Australasia      22  
         Australia       923  
         Austria         569  
         Azerbaijan       43  
         Bahamas         24  
         Bahrain          1  
         Barbados         1  
         Belarus         154  
         Belgium         276  
         Bermuda          1  
         Bohemia          5  
         Botswana         2  
         Brazil          184  
         British West Indies 2  
         Bulgaria        411  
         Burundi          3  
         Cameroon        12  
         Canada          846  
         Chile           24  
         China          1120  
         Colombia        29  
         Costa Rica        7  
         Ivory Coast       2  
         Croatia          67  
         Cuba            420  
         Cyprus           2  
         ...  
         Spain           268  
         Sri Lanka         4  
         Sudan             2  
         Suriname          4
```

Sweden	1217
Switzerland	630
Syria	6
Chinese Taipei	32
Tajikistan	4
Tanzania	4
Thailand	44
Togo	1
Tonga	2
Trinidad and Tobago	27
Tunisia	19
Turkey	191
Uganda	14
Ukraine	220
United Arab Emirates	3
United States	5684
Uruguay	16
Uzbekistan	38
Venezuela	18
Vietnam	4
Virgin Islands	2
Yugoslavia	171
Independent Olympic Participants	4
Zambia	3
Zimbabwe	18
Mixed team	38

Name: Total Points, dtype: int64

## 1.1 Part 2

For the next set of questions, we will be using census data from the [United States Census Bureau](#). Counties are political and geographic subdivisions of states in the United States. This dataset contains population data for counties and states in the US from 2010 to 2015. [See this document](#) for a description of the variable names.

The census dataset (census.csv) should be loaded as census\_df. Answer questions using this as appropriate.

### 1.1.1 Question 5

Which state has the most counties in it? (hint: consider the sumlevel key carefully! You'll need this for future questions too...)

*This function should return a single string value.*

```
In [15]: census_df = pd.read_csv('census.csv')
        census_df.head(10
        )
```

```
Out[15]:  SUMLEV  REGION  DIVISION  STATE  COUNTY  STNAME  CTYNAME \
         0         40         3         6         1         0  Alabama  Alabama
```

1	50	3	6	1	1	Alabama	Autauga County
2	50	3	6	1	3	Alabama	Baldwin County
3	50	3	6	1	5	Alabama	Barbour County
4	50	3	6	1	7	Alabama	Bibb County
5	50	3	6	1	9	Alabama	Blount County
6	50	3	6	1	11	Alabama	Bullock County
7	50	3	6	1	13	Alabama	Butler County
8	50	3	6	1	15	Alabama	Calhoun County
9	50	3	6	1	17	Alabama	Chambers County

	CENSUS2010POP	ESTIMATESBASE2010	POPESTIMATE2010	...	\
0	4779736	4780127	4785161	...	
1	54571	54571	54660	...	
2	182265	182265	183193	...	
3	27457	27457	27341	...	
4	22915	22919	22861	...	
5	57322	57322	57373	...	
6	10914	10915	10887	...	
7	20947	20946	20944	...	
8	118572	118586	118437	...	
9	34215	34170	34098	...	

	RDOMESTICMIG2011	RDOMESTICMIG2012	RDOMESTICMIG2013	RDOMESTICMIG2014	\
0	0.002295	-0.193196	0.381066	0.582002	
1	7.242091	-2.915927	-3.012349	2.265971	
2	14.832960	17.647293	21.845705	19.243287	
3	-4.728132	-2.500690	-7.056824	-3.904217	
4	-5.527043	-5.068871	-6.201001	-0.177537	
5	1.807375	-1.177622	-1.748766	-2.062535	
6	-30.953709	-5.180127	-1.130263	14.354290	
7	-14.032727	-11.684234	-5.655413	1.085428	
8	-6.155670	-4.611706	-5.524649	-4.463211	
9	-2.731639	3.849092	2.872721	-2.287222	

	RDOMESTICMIG2015	RNETMIG2011	RNETMIG2012	RNETMIG2013	RNETMIG2014	\
0	-0.467369	1.030015	0.826644	1.383282	1.724718	
1	-2.530799	7.606016	-2.626146	-2.722002	2.592270	
2	17.197872	15.844176	18.559627	22.727626	20.317142	
3	-10.543299	-4.874741	-2.758113	-7.167664	-3.978583	
4	0.177258	-5.088389	-4.363636	-5.403729	0.754533	
5	-1.369970	1.859511	-0.848580	-1.402476	-1.577232	
6	-16.167247	-29.001673	-2.825524	1.507017	17.243790	
7	-6.529805	-13.936612	-11.586865	-5.557058	1.184103	
8	-3.376322	-5.791579	-4.092677	-5.062836	-3.912834	
9	1.349468	-1.821092	4.701181	3.781439	-1.290228	

	RNETMIG2015
0	0.712594

```

1    -2.187333
2    18.293499
3   -10.543299
4     1.107861
5    -0.884411
6   -13.193961
7    -6.430868
8    -2.806406
9     2.346901

```

```
[10 rows x 100 columns]
```

```

In [17]: def answer_five():
          new_census_df = census_df[census_df['SUMLEV']== 50]
          return new_census_df.groupby('STNAME').count()['SUMLEV'].idxmax()

          answer_five()

```

```
Out[17]: 'Texas'
```

### 1.1.2 Question 6

**Only looking at the three most populous counties for each state**, what are the three most populous states (in order of highest population to lowest population)? Use CENSUS2010POP.

*This function should return a list of string values.*

```

In [26]: def answer_six():
          new_census_df = census_df[census_df['SUMLEV']== 50]
          high_pop = new_census_df.sort_values('CENSUS2010POP', ascending = False).groupby('STNAME').sum()
          return high_pop.groupby('STNAME').sum().sort_values('CENSUS2010POP', ascending = False)

          answer_six()

```

```
Out[26]: ['California', 'Texas', 'Illinois']
```

### 1.1.3 Question 7

Which county has had the largest absolute change in population within the period 2010-2015? (Hint: population values are stored in columns POPESTIMATE2010 through POPESTIMATE2015, you need to consider all six columns.)

e.g. If County Population in the 5 year period is 100, 120, 80, 105, 100, 130, then its largest change in the period would be  $|130-80| = 50$ .

*This function should return a single string value.*

```

In [34]: def answer_seven():
          new_census_df = census_df[census_df['SUMLEV']== 50][['CTYNAME', 'POPESTIMATE2010', 'POPESTIMATE2011', 'POPESTIMATE2012', 'POPESTIMATE2013', 'POPESTIMATE2014', 'POPESTIMATE2015']]
          new_census_df['max_diff'] = abs(new_census_df.max(axis=1)-new_census_df.min(axis=1))
          max_val = new_census_df.sort_values('max_diff', ascending= False)
          return max_val.iloc[0][0]

          answer_seven()

```

```
Out[34]: 'Harris County'
```

#### 1.1.4 Question 8

In this datafile, the United States is broken up into four regions using the "REGION" column.

Create a query that finds the counties that belong to regions 1 or 2, whose name starts with 'Washington', and whose POPESTIMATE2015 was greater than their POPESTIMATE 2014.

*This function should return a 5x2 DataFrame with the columns = ['STNAME', 'CTYNAME'] and the same index ID as the census\_df (sorted ascending by index).*

```
In [38]: def answer_eight():
         new_census_df = census_df[census_df['SUMLEV']== 50]
         pop_new = new_census_df[((new_census_df['REGION']==1)|(new_census_df['REGION']==2))
         return pop_new
```

```
answer_eight()
```

```
Out[38]:
```

	STNAME	CTYNAME
896	Iowa	Washington County
1419	Minnesota	Washington County
2345	Pennsylvania	Washington County
2355	Rhode Island	Washington County
3163	Wisconsin	Washington County

```
In [ ]:
```